

In [31]:

```

import json
import pandas as pd

#Import Required Module
#!pip install tensorflow
#!pip install keras
#!pip install wordcloud
#!pip install libomp
#!pip install --upgrade libomp
#!pip3 install xgboost
#!pip install --upgrade xgboost #ran this command on terminal on mac OS 'conda i
#!pip install gensim
#!pip install pronouncing
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import re,string
import nltk
from nltk import word_tokenize
from nltk.corpus import stopwords
from nltk.probability import FreqDist
stops = set(stopwords.words("english"))
punctuation = string.punctuation
from textblob import TextBlob
from sklearn.feature_extraction.text import CountVectorizer,TfidfVectorizer
from sklearn import model_selection, preprocessing, linear_model, naive_bayes, m
from sklearn.metrics import classification_report,roc_curve,confusion_matrix, auc
from sklearn.tree import DecisionTreeClassifier,ExtraTreeClassifier
from sklearn.ensemble import RandomForestClassifier,ExtraTreesClassifier,AdaBoos
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import label_binarize
from sklearn.naive_bayes import MultinomialNB
from sklearn.multiclass import OneVsRestClassifier
from sklearn.linear_model import LogisticRegression
from keras import layers, models, optimizers
from sklearn.decomposition import PCA,TruncatedSVD
import xgboost
import gensim
import warnings
warnings.filterwarnings('ignore')
import pronouncing

```

In [32]:

```

df_business = pd.read_json('yelp_training_set_business.json', lines=True)
print(df_business.head())

df_business.shape

```

	business_id	full_address \
0	rncjoVoEFUJGCUoC1JgnUA	8466 W Peoria Ave\nSte 6\nPeoria, AZ 85345
1	0FNFSzCFP_rGUoJx8W7tJg	2149 W Wood Dr\nPhoenix, AZ 85029
2	3f_lyB6vFK48ukH6ScvLHg	1134 N Central Ave\nPhoenix, AZ 85004

```

3 usAsSV36QmUej8--yvN-dg      845 W Southern Ave\nPhoenix, AZ 85041
4 PzOqRohWw7F7YEPBz6AubA 6520 W Happy Valley Rd\nSte 101\nGlendale Az, ...

```

```

      open      categories      city \
0 True  [Accountants, Professional Services, Tax Servi...  Peoria
1 True      [Sporting Goods, Bikes, Shopping]  Phoenix
2 True      []  Phoenix
3 True      [Food, Grocery]  Phoenix
4 True      [Food, Bagels, Delis, Restaurants]  Glendale Az

```

```

      review_count      name neighborhoods      longitude state \
0          3      Peoria Income Tax Service      [] -112.241596  AZ
1          5      Bike Doctor      [] -112.105933  AZ
2          4  Valley Permaculture Alliance      [] -112.073933  AZ
3          5      Food City      [] -112.085377  AZ
4         14      Hot Bagels & Deli      [] -112.200264  AZ

```

```

      stars      latitude      type
0      5.0  33.581867  business
1      5.0  33.604054  business
2      5.0  33.460526  business
3      3.5  33.392210  business
4      3.5  33.712797  business

```

Out[32]: (11537, 13)

In [33]:

```

import json
import pandas as pd

from ast import literal_eval

json_read = pd.read_json('yelp_training_set_review.json', orient="records",line

df_review = pd.concat(json_read)
print(df_review.head())
df_review.shape

df = pd.DataFrame(df_review['votes'].values.tolist(), index=df_review.index)
print (df)

```

```

      votes      user_id \
0 {'funny': 0, 'useful': 5, 'cool': 2} rLt18ZkDX5vH5nAx9C3q5Q
1 {'funny': 0, 'useful': 0, 'cool': 0} 0a2KyEL0d3Yb1V6aivbIuQ
2 {'funny': 0, 'useful': 1, 'cool': 0} 0hT2KtFLiobPvh6cDC8JQg
3 {'funny': 0, 'useful': 2, 'cool': 1} uZet19T0NcROGOyFfughhg
4 {'funny': 0, 'useful': 0, 'cool': 0} vYmM4KTsC8ZfQBg-j5MWkw

```

```

      review_id      stars      date \
0 fWKvX83p0-ka4JS3dc6E5A      5 2011-01-26
1 IjZ33sJrzXqU-0X6U8NwyA      5 2011-07-27
2 IESLBzqUCLdSzSqm0eCSxQ      4 2012-06-14
3 G-WvGaISbqqqAMHlNnByoda      5 2010-05-27
4 1uJFq2r5QfJG_6ExMRCaGw      5 2012-01-05

```

```

      text      type \
0 My wife took me here on my birthday for breakf...  review
1 I have no idea why some people give bad review...  review
2 love the gyro plate. Rice is so good and I als...  review
3 Rosie, Dakota, and I LOVE Chaparral Dog Park!!...  review
4 General Manager Scott Petello is a good egg!!!!  review

```

```

      business_id
0  9yKzy9PApeiPPOUJEtnvkg
1  ZRJwVLyzEJq1VAihDhYiow
2  6oRAC4uyJCsj11X0WZpVSA
3  _1QQZuf4zZOyFCvXc0o6Vg
4  6ozycU1RpktNG2-1BroVtw
      funny  useful  cool
0         0      5     2
1         0      0     0
2         0      1     0
3         0      2     1
4         0      0     0
...      ...      ...   ...
229902     0      0     0
229903     0      2     0
229904     0      0     0
229905     1      2     0
229906     1      0     1

```

[229907 rows x 3 columns]

In [34]:

```

df_user = pd.read_json('yelp_training_set_user.json', lines=True)
print(df_user.head())

df_user.shape

```

```

      votes      user_id      name \
0  {'funny': 0, 'useful': 7, 'cool': 0} CR2y7yEm4X035ZMzrTtN9Q      Jim
1  {'funny': 0, 'useful': 1, 'cool': 0} _9GXoHhdx30ujPaQwh6Ew      Kelle
2  {'funny': 0, 'useful': 1, 'cool': 0} 8mM-nqxjg6pT04kwcjMbsw  Stephanie
3  {'funny': 0, 'useful': 2, 'cool': 0} Ch6CdTR2IVaVAnr-RglMOg      T
4  {'funny': 0, 'useful': 0, 'cool': 0} NZrLmHRYiHmyTlJrfzkCOA      Beth

      average_stars  review_count  type
0         5.0             6  user
1         1.0             2  user
2         5.0             2  user
3         5.0             2  user
4         1.0             1  user

```

Out[34]: (43873, 6)

In [35]:

```

df_checkin = pd.read_json('yelp_training_set_checkin.json', lines=True)
print(df_checkin.head())

```

```

      checkin_info      type \
0  {'11-3': 17, '8-5': 1, '15-0': 2, '15-3': 2, '...' checkin
1  {'0-5': 1, '2-6': 2, '2-5': 3, '3-6': 1, '3-5'... checkin
2  {'13-4': 1, '7-4': 1, '15-3': 1, '18-5': 1, '2... checkin
3  {'13-5': 1, '17-6': 1, '15-1': 1, '20-0': 1, '...' checkin
4  {'16-2': 1, '14-5': 1, '12-5': 2, '15-4': 1, '...' checkin

      business_id
0  KO9CpaSP0oqm0iCWm5scmg
1  oRqBAYtcBYZHXa7G8FlPaA
2  6cy2C9aBXUwkrh4bY1DApw
3  D0IB17N66FiyYDCzTlAI4A
4  HLQGo3EaYVvAv22bONGkIw

```

In [36]:

```

df_merge1=df_review.merge(df_business,how='left', on='business_id')

```

```
df_merge1.head()
```

Out[36]:

	votes	user_id	review_id	stars_x	date	text	type_
0	{'funny': 0, 'useful': 5, 'cool': 2}	rLtl8ZkDX5vH5nAx9C3q5Q	fWKvX83p0-ka4JS3dc6E5A	5	2011-01-26	My wife took me here on my birthday for breakf...	revie
1	{'funny': 0, 'useful': 0, 'cool': 0}	0a2KyEL0d3Yb1V6aivbluQ	ljZ33sJrzXqU-0X6U8NwyA	5	2011-07-27	I have no idea why some people give bad review...	revie
2	{'funny': 0, 'useful': 1, 'cool': 0}	0hT2KtLiobPvh6cDC8JQg	IESLBzqUCLdSzSqm0eCSxQ	4	2012-06-14	love the gyro plate. Rice is so good and I als...	revie
3	{'funny': 0, 'useful': 2, 'cool': 1}	uZetl9T0NcROGOyFfughhg	G-WvGalSbqqaMHINnByodA	5	2010-05-27	Rosie, Dakota, and I LOVE Chaparral Dog Park!!...	revie
4	{'funny': 0, 'useful': 0, 'cool': 0}	vYmM4KTsC8ZfQBg-j5MWkw	1uJFq2r5QfJG_6ExMRCaGw	5	2012-01-05	General Manager Scott Petello is a good egg!!!...	revie

In [37]:

```
df_merge2=df_merge1.merge(df_checkin,how='left', on='business_id')
df_merge2.head()
```

Out[37]:

	votes	user_id	review_id	stars_x	date	text	type_
0	{'funny': 0, 'useful': 5, 'cool': 2}	rLtl8ZkDX5vH5nAx9C3q5Q	fWKvX83p0-ka4JS3dc6E5A	5	2011-01-26	My wife took me here on my birthday for breakf...	revie
1	{'funny': 0, 'useful': 0, 'cool': 0}	0a2KyEL0d3Yb1V6aivbluQ	ljZ33sJrzXqU-0X6U8NwyA	5	2011-07-27	I have no idea why some people give bad review...	revie

	votes	user_id	review_id	stars_x	date	text	type_
2	{'funny': 0, 'useful': 1, 'cool': 0}	0hT2KtfLiobPvh6cDC8JQg	IESLBzqUCLdSzSqm0eCSxQ	4	2012-06-14	love the gyro plate. Rice is so good and I als...	revie
3	{'funny': 0, 'useful': 2, 'cool': 1}	uZetl9T0NcROGOyFfughhg	G-WvGalSbqqaMHINnByodA	5	2010-05-27	Rosie, Dakota, and I LOVE Chaparral Dog Park!!...	revie
4	{'funny': 0, 'useful': 0, 'cool': 0}	vYmM4KTsC8ZfQBg-j5MWkw	1uJFq2r5QfJG_6ExMRCaGw	5	2012-01-05	General Manager Scott Petello is a good egg!!!...	revie

5 rows × 22 columns

```
In [38]: df_merge4=df_merge2.merge(df_user,how='left', on='user_id')
df_merge4.head()
```

	votes_x	user_id	review_id	stars_x	date	text	type_
0	{'funny': 0, 'useful': 5, 'cool': 2}	rLtl8ZkDX5vH5nAx9C3q5Q	fWKvX83p0-ka4JS3dc6E5A	5	2011-01-26	My wife took me here on my birthday for breakf...	revie
1	{'funny': 0, 'useful': 0, 'cool': 0}	0a2KyEL0d3Yb1V6aivbluQ	ljZ33sJrzXqU-0X6U8NwyA	5	2011-07-27	I have no idea why some people give bad review...	revie
2	{'funny': 0, 'useful': 1, 'cool': 0}	0hT2KtfLiobPvh6cDC8JQg	IESLBzqUCLdSzSqm0eCSxQ	4	2012-06-14	love the gyro plate. Rice is so good and I als...	revie
3	{'funny': 0, 'useful': 2, 'cool': 1}	uZetl9T0NcROGOyFfughhg	G-WvGalSbqqaMHINnByodA	5	2010-05-27	Rosie, Dakota, and I LOVE Chaparral Dog Park!!...	revie

	votes_x	user_id	review_id	stars_x	date	text	type_
	{'funny': 0, 'useful': 0, 'cool': 0}	vYmM4KTsC8ZfQBg- j5MWkw	1uJFq2r5QfJG_6ExMRCaGw	5	2012- 01-05	General Manager Scott Petello is a good egg!!!!...	revie

5 rows × 27 columns

```
In [39]: df_merge3=df_merge4.head(50000)
df_merge3.shape
```

Out[39]: (50000, 27)

```
In [19]: #Step6: Plot Word cloud for 1 star rating restaurants
print('\nWord cloud for 1 star rating restaurants\n')
from subprocess import check_output
from wordcloud import WordCloud, STOPWORDS

stopwords = set(STOPWORDS)
wordcloud = WordCloud(
    background_color='black',
    stopwords=stopwords,
    max_words=200,
    max_font_size=40,
    random_state=42
).generate(str(df_merge3[df_merge3['stars_x']==1]['name

fig = plt.figure(1,figsize=(12,18))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()

#Step6: Plot Word cloud for 2 star rating restaurants
print('\nWord cloud for 2 star rating restaurants\n')
from subprocess import check_output
from wordcloud import WordCloud, STOPWORDS

stopwords = set(STOPWORDS)
wordcloud = WordCloud(
    background_color='black',
    stopwords=stopwords,
    max_words=200,
    max_font_size=40,
    random_state=42
).generate(str(df_merge3[df_merge3['stars_x']==2]['name

fig = plt.figure(1,figsize=(12,18))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()

#Step6: Plot Word cloud for 3 star rating restaurants
```

```

print('\nWord cloud for 3 star rating restaurants\n')
from subprocess import check_output
from wordcloud import WordCloud, STOPWORDS

stopwords = set(STOPWORDS)
wordcloud = WordCloud(
    background_color='black',
    stopwords=stopwords,
    max_words=200,
    max_font_size=40,
    random_state=42
).generate(str(df_merge3[df_merge3['stars_x']==3]['name

fig = plt.figure(1,figsize=(12,18))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()

#Step6: Plot Word cloud for 4 star rating restaurants
print('\nWord cloud for 4 star rating restaurants\n')
from subprocess import check_output
from wordcloud import WordCloud, STOPWORDS

stopwords = set(STOPWORDS)
wordcloud = WordCloud(
    background_color='black',
    stopwords=stopwords,
    max_words=200,
    max_font_size=40,
    random_state=42
).generate(str(df_merge3[df_merge3['stars_x']==4]['name

fig = plt.figure(1,figsize=(12,18))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()


#Step6: Plot Word cloud for 5 star rating restaurants
print('\nWord cloud for 5 star rating restaurants\n')
from subprocess import check_output
from wordcloud import WordCloud, STOPWORDS

stopwords = set(STOPWORDS)
wordcloud = WordCloud(
    background_color='black',
    stopwords=stopwords,
    max_words=200,
    max_font_size=40,
    random_state=42
).generate(str(df_merge3[df_merge3['stars_x']==5]['name


fig = plt.figure(1,figsize=(12,18))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()

```

Word cloud for 1 star rating restaurants



Word cloud for 2 star rating restaurants




Word cloud for 3 star rating restaurants



Word cloud for 4 star rating restaurants



Word cloud for 5 star rating restaurants




In [20]:

```
##### VISUALIZATION #####

#Step4: Create a pie chart to show the percentage wise category distribution
print('\nPie Chart:\n')
labels = '1', '2', '3', '4', '5'
sizes = [17516, 20957, 35363, 79878, 76193]
fig1, ax1 = plt.subplots(figsize=(5,5))
ax1.pie(sizes, labels=labels, autopct='%1.1f%%',
        shadow=False, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.show()
```

Pie Chart:



In [29]:

```
import matplotlib.pyplot as plt
num_bins = 5
```

```
fig, ax = plt.subplots(figsize = (10,7))

n, bins, patches = ax.hist(df_merge3['stars_x'], num_bins, facecolor='#2b8cbe',



ax.set_title('Histogram of Ratings', fontsize = 15, pad=15)
ax.set_xlabel('rating')
ax.set_ylabel('frequency')

plt.show()


import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

x = df_merge3['stars_x']
y = df_merge3['stars_x'].value_counts(ascending=True)
fig, ax = plt.subplots(figsize=(12,10) )
width = 0.75 # the width of the bars
df_merge3['stars_x'].value_counts().plot(kind='bar');
plt.title("Star Rating Distribution")
plt.ylabel('# of businesses', fontsize=12)
plt.xlabel('Star Ratings ', fontsize=12)
plt.show()

#Step4: Create a pie chart to show the percentage wise rating distribution
print('\nPie Chart:\n')
labels = '1', '2', '3', '4', '5'
sizes = [17516, 20957, 35363, 79878, 76193]
fig1, ax1 = plt.subplots(figsize=(5,5))
ax1.pie(sizes, labels=labels, autopct='%1.1f%%',
        shadow=False, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.show()
```



Pie Chart:



In [40]:

```
# bin the data into negative, neutral, and positive values
bins = [0, 2, 4, 6]

bin_names = ['negative', 'neutral', 'positive']

score_bin = pd.Series(df_merge3.stars_x, name = 'score')

score = pd.cut(score_bin, bins, labels=bin_names, right=False)

# number of counts per score

pd.value_counts(score)

df_merge3 = pd.concat([df_merge3, score], axis=1)

df_merge3.head(2)
```

Out[40]:

	votes_x	user_id	review_id	stars_x	date	text	type_x
0	{'funny': 0, 'useful': 5, 'cool': 2}	rLtI8ZkDX5vH5nAx9C3q5Q	fWKvX83p0-ka4JS3dc6E5A	5	2011-01-26	My wife took me here on my birthday for breakf...	review
1	{'funny': 0, 'useful': 0, 'cool': 0}	0a2KyEL0d3Yb1V6aivbIuQ	IjZ33sJrzXqU-0X6U8NwyA	5	2011-07-27	I have no idea why some people give bad review...	review

2 rows x 28 columns

In [41]:

```

# number of counts per score

top_restaurants=df_merge3[(df_merge3['stars_x']==5 ) ]

# top 10 restaurants with most reviews
#top_restaurants_10= top_restaurants.head(10)
top_restaurants=top_restaurants.drop_duplicates(subset=['latitude','longitude','
top_restaurants_10 = top_restaurants.sort_values(['review_count_x'], ascending=[
print(top_restaurants_10.head())
#!pip install folium pandas
import folium
#!pip install --upgrade pandas

top_restaurants_10=top_restaurants_10[['latitude','longitude','name_x','review_c

#top_restaurants_10 = top_restaurants_10.sort_values(['review_count'], ascending

center = [33.581867      , -112.241596]
map_USA = folium.Map(location=center, zoom_start=8)
for index, top_restaurants_10 in top_restaurants_10.iterrows():
    location = [top_restaurants_10['latitude'], top_restaurants_10['longitude']]
    folium.Marker(location, popup = f'Latitude:{top_restaurants_10["latitude"]}\

map_USA

```

		votes_x	user_id \
2493	{'funny': 1, 'useful': 2, 'cool': 1}	-vRFUY8ixuNniCCNVvmkRQ	
205	{'funny': 3, 'useful': 2, 'cool': 2}	qklF6QU-bi4Y4Bt4g6Sv8A	
1321	{'funny': 0, 'useful': 0, 'cool': 0}	TL46g36OKxmgSDYRFJdPRg	
1791	{'funny': 0, 'useful': 2, 'cool': 2}	8-2W5CmkDl9vrkxRpkIPRg	
600	{'funny': 2, 'useful': 3, 'cool': 3}	5RxpP2Woo7CpOGUmKgDyAw	

	review_id	stars_x	date \
2493	7VsG6-m3wNhuoMUeHQ8Vcw	5	2012-06-03
205	h2c-jLW9cLViiiubGqTaog	5	2011-01-25
1321	XdAtb3hIzydAmzdFy0Ubyw	5	2010-08-05
1791	3eJ2M94guIOhVWexTM6iEw	5	2010-10-31
600	R3Dndbcc0jEAlXyr4JV1EA	5	2008-05-08

	text	type_x \
2493	It's called America's friendliest airport for ...	review
205	On one of my many visits to see mi amore, he t...	review
1321	Still solid, still delicious. I love this place!	review
1791	Matt's was absolutely fantastic. Got there at ...	review
600	After reading the great reviews and getting pl...	review

	business_id	full_address	open \
2493	hW0Ne_HTHEAgGF1rAdmR-g	3400 E Sky Harbor Blvd\nPhoenix, AZ 85034	True
205	VVeogjZya58oiTxK7qUjAQ	623 E Adams St\nPhoenix, AZ 85004	True
1321	JokKtdXU7zXHcr20Lrk29A	1340 E 8th St\nSte 104\nTempe, AZ 85281	True
1791	ntN85eu27C04nwyPa8Ihtw	801 N 1st St\nPhoenix, AZ 85004	True
600	EWMwV5V9BxNs_U6nNVMeqw	3815 N Central Ave\nPhoenix, AZ 85012	True

	latitude	type_y \
2493	33.434750	business

```

205    ...    33.449233    business
1321   ...    33.419451    business
1791   ...    33.456696    business
600    ...    33.491645    business

```

```

                                checkin_info    type_x  \
2493  {'2-5': 1, '22-6': 192, '22-5': 42, '22-4': 92...  checkin
205    {'10-0': 2, '17-0': 6, '18-4': 19, '13-5': 9, ...  checkin
1321   {'22-6': 11, '22-5': 36, '22-4': 41, '22-3': 2...  checkin
1791   {'7-0': 7, '11-3': 18, '3-0': 1, '5-6': 5, '11...  checkin
600    {'22-6': 5, '22-5': 10, '22-4': 12, '22-3': 2,...  checkin

```

```

                                votes_y    name_y  average_stars  \
2493          {'funny': 49, 'useful': 104, 'cool': 51}  Brittany          3.93
205           {'funny': 11, 'useful': 21, 'cool': 9}    Becca          3.89
1321  {'funny': 1050, 'useful': 1809, 'cool': 1331}  Jennifer          3.74
1791           {'funny': 49, 'useful': 124, 'cool': 51}    Chris          3.87
600          {'funny': 154, 'useful': 382, 'cool': 266}    Emily          4.19

```

```

        review_count_y  type_y    score
2493          58.0    user  positive
205          28.0    user  positive
1321        650.0    user  positive
1791          83.0    user  positive
600        169.0    user  positive

```

```
[5 rows x 28 columns]
```

Out[41]: Make this Notebook Trusted to load map: File -> Trust Notebook

In [45]:

```

# number of counts per score

topRatedRestaurants = pd.Series(df_merge3['name_x'])
topRestaurantsCounts = pd.value_counts(topRatedRestaurants)

# top 10 restaurants with most reviews
topRestaurantsCounts.head(10)

```

```

Out[45]: Pita Jungle          273
         Oregano's Pizza Bistro  199

```

Cornish Pasty Company	192
Phoenix Sky Harbor International Airport	186
Lo-Lo's Chicken & Waffles	182
Pizzeria Bianco	177
Four Peaks Brewing Co	159
Matt's Big Breakfast	158
Postino Arcadia	144
Cibo	144

Name: name_x, dtype: int64

In [46]:

```
import matplotlib.pyplot as plt

df_merge3['text_cleaned'] = df_merge3['text'].apply(lambda x: x.split())
df_merge3.head()

from collections import defaultdict
word_freq = defaultdict(int)
for sent in df_merge3['text_cleaned']:
    for i in sent:
        word_freq[i] += 1
len(word_freq)

sorted(word_freq, key=word_freq.get, reverse=True)[:10]

sentences = df_merge3['text_cleaned']
# Set values for various parameters
num_features = 100      # Word vector dimensionality
min_word_count = 40     # ignore all words with total frequency lower than this
num_workers = 4         # Number of threads to run in parallel
context = 10            # Context window size

# Initialize and train the model (this will take some time)
from gensim.models import word2vec

model = word2vec.Word2Vec(sentences,
                          workers=num_workers,
                          size=num_features,
                          min_count=min_word_count,
                          window=context)

print("Training finished!")

# save the model for later use. You can load it later using Word2Vec.load()
model_name = "Word_Embedding"
model.save(model_name)

# Get vocabulary count of the model
```



```
vocab_tmp = list(model.wv.vocab)
print('Vocab length:', len(vocab_tmp))

from sklearn.metrics.pairwise import cosine_similarity

model.similarity('dish', 'plate')

model.most_similar(positive=['tasty', 'pleased', 'health', 'enjoy'], negative=['ba

from gensim.models import Word2Vec
# Load the trained modelNumeric Representations of Words
model = Word2Vec.load("Word_Embedding")

vocab_tmp = list(model.wv.vocab)
print('Vocab length:', len(vocab_tmp))

# Get distributional representation of each word
X = model[vocab_tmp]

from sklearn import decomposition
# get two principle components of the feature space
pca = decomposition.PCA(n_components=2).fit_transform(X)

good_list = [x for x,y in model.most_similar('great', topn=20)]
bad_list = [x for x,y in model.most_similar('worst', topn=20)]
# good_list = [x for x,y in model.most_similar(positive=['good', 'great', 'health

# set figure settings
plt.figure(figsize=(15,15))

# save pca values and vocab in dataframe df
df = pd.concat([pd.DataFrame(pca), pd.Series(vocab_tmp)], axis=1)
df.columns = ['x', 'y', 'word']

plt.xlabel("1st principal component", fontsize=14)
plt.ylabel("2nd principal component", fontsize=14)

plt.scatter(x=df['x'], y=df['y'], s=3, alpha=0.3)

good_words = df[df['word'].isin(good_list)]['word']
for i, word in good_words.items():
    plt.annotate(word, (df['x'].iloc[i], df['y'].iloc[i]), fontsize=16, color='gre

bad_words = df[df['word'].isin(bad_list)]['word']
for i, word in bad_words.items():
    plt.annotate(word, (df['x'].iloc[i], df['y'].iloc[i]), fontsize=13, color='red
```



```

df_merge3['upper_case_word_count'] = df_merge3['text'].apply(lambda x: len([wrds for wrds in x.split() if wrds.isupper()]))
df_merge3['stopword_count'] = df_merge3['text'].apply(lambda x: len([wrds for wrds in x.split() if wrds in stopwords]))

df_merge3['line_count'] = df_merge3['text'].apply(lambda x: len([line for line in x.splitlines() if line]))
#print(df['line_count'])
print('\nPrint NLP/Text based features:\n')
print(df_merge3[['char_count', 'word_count', 'word_density', 'punctuation_count', 'title_word_count', 'upper_case_word_count', 'stopword_count', 'line_count']])

```

Print NLP/Text based features:

	char_count	word_count	word_density	punctuation_count	title_word_count	\
0	889	155	5.698718	21	16	
1	1345	257	5.213178	36	25	
2	76	16	4.470588	3	2	
3	419	76	5.441558	18	13	
4	469	86	5.390805	38	12	
5	2094	366	5.705722	64	33	
6	1565	292	5.341297	50	45	
7	274	50	5.372549	9	9	
8	349	62	5.539683	13	9	
9	186	34	5.314286	4	4	

	upper_case_word_count	stopword_count	line_count
0	3	71	1
1	6	134	1
2	1	6	1
3	2	33	1
4	2	44	1
5	10	171	1
6	11	134	1
7	3	25	1
8	1	29	1
9	0	16	1

In [49]:

```

counts_df = df_merge3[['score', 'text', 'word_count']]

# separate by positive and negative reviews
counts_pos = counts_df.loc[(counts_df['score']=='positive')]
counts_neg = counts_df.loc[(counts_df['score']=='negative')]

# create figure
fig, ax = plt.subplots(figsize = (12,10))

sns.boxplot(x=counts_df['score'], y=counts_df['word_count'])

# title
ax.set_title('Number of Words in the Reviews', fontsize = 15, loc = 'left')

# set x axis label
ax.set_xlabel('Sentiment of Review', fontsize = 13)

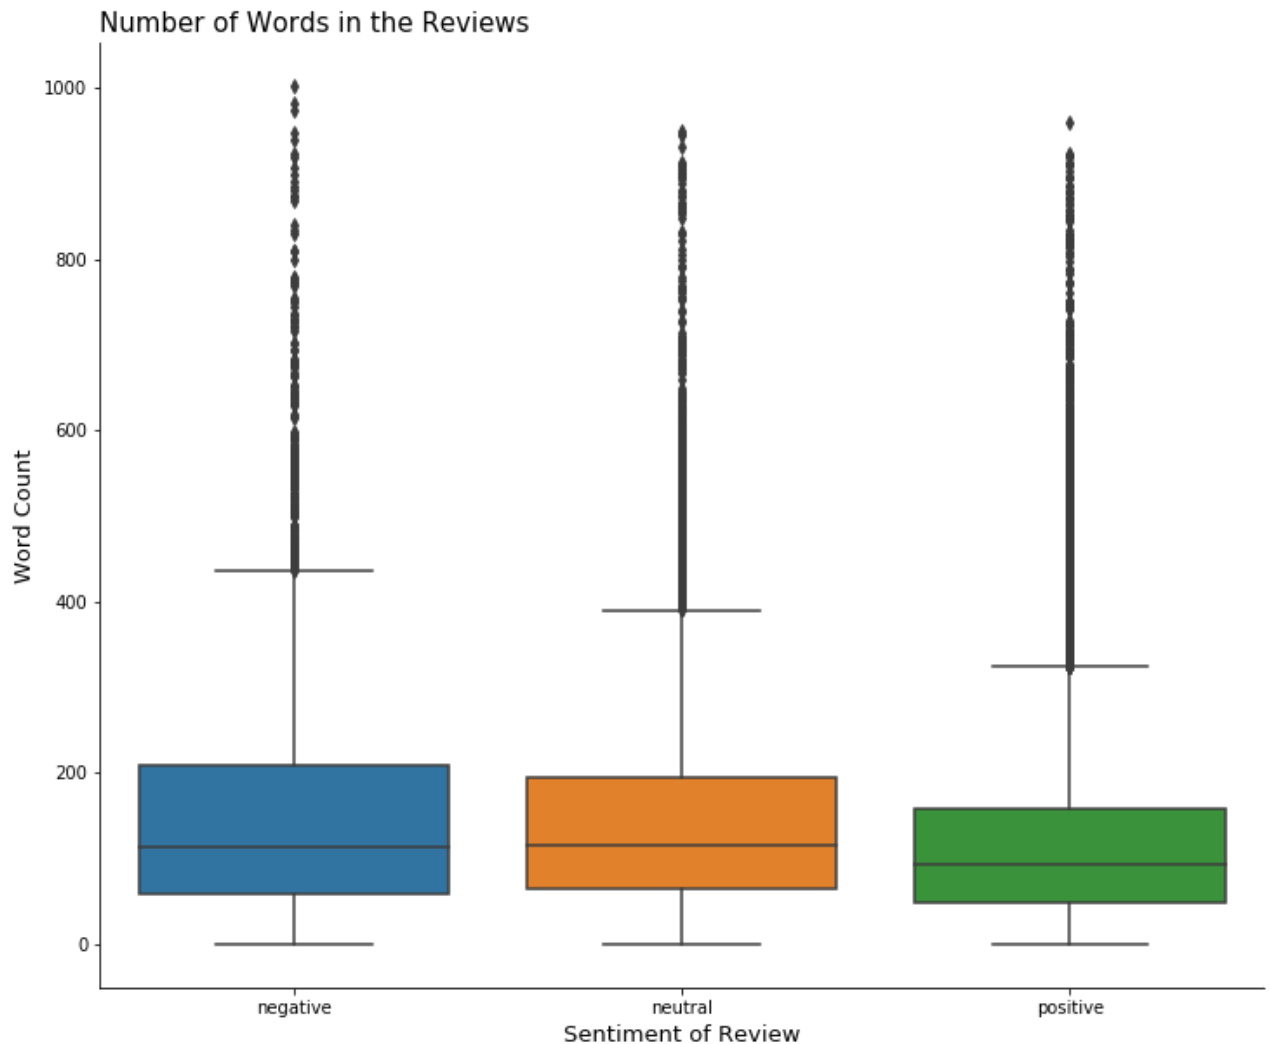
# set y axis label
ax.set_ylabel('Word Count', fontsize = 13)

# remove spines
ax.spines['top'].set_visible(False)
ax.spines['right'].set_visible(False)
ax.spines['left'].set_visible(True)

```

```
ax.spines['bottom'].set_visible(True)

plt.show()
```



In [51]: *#Step 12: Clean text: no punctuation/all lowercase/remove stop words*

```
#Convert all cases to lower case
df_merge3 = df_merge3.astype(str).apply(lambda x: x.str.lower())
#print('\nFew sample records after converting strings to low case:\n')
#print(df.head())

#Remove the punctuations from the dataframe
def remove_punctuations(text):
    for punctuation in string.punctuation:
        text = text.replace(punctuation, '')
    return text

df_merge3['text'] = df_merge3['text'].apply(remove_punctuations)
#print('\nFew sample records after removing punctuations:\n')
#print(df.head())
```

```
#Remove stop words from dataframe
df_merge3['text'] = df_merge3['text'].apply(lambda x: ' '.join([item for item in x if item not in STOP_WORDS]))
#print('\nFew sample records after removing stop words:\n')
#print(df.head())

#Apply porter_stemmer on dataframe
from nltk.stem.porter import PorterStemmer
porter_stemmer = PorterStemmer()

df_merge3['txt_tokenized']=df_merge3['text'].apply(lambda x : filter(None,x.split()))
df_merge3['txt_stemmed']=df_merge3['txt_tokenized'].apply(lambda x : [porter_stemmer.stem(word) for word in x])
df_merge3['txt_stemmed_sentence']=df_merge3['txt_stemmed'].apply(lambda x : " ".join(x))
print('\nFew sample records after doing cleaning/preprocessing (convert to low case/remove punctuation/remove stopwords/apply porter stemmer):\n')
print(df_merge3.head())
```

Few sample records after doing cleaning/preprocessing (convert to low case/remove punctuation/remove stopwords/apply porter stemmer):

	votes_x	user_id \
0	{'funny': 0, 'useful': 5, 'cool': 2}	rltl8zkdx5vh5nax9c3q5q
1	{'funny': 0, 'useful': 0, 'cool': 0}	0a2kyel0d3yblv6aivbiuq
2	{'funny': 0, 'useful': 1, 'cool': 0}	0ht2ktfliobpvh6cdc8jqg
3	{'funny': 0, 'useful': 2, 'cool': 1}	uzetl9t0ncrogoyffughhg
4	{'funny': 0, 'useful': 0, 'cool': 0}	vymm4ktsc8zfqbqg-j5mwkw

	review_id	stars_x	date \
0	fwkvx83p0-ka4js3dc6e5a	5	2011-01-26
1	ijz33sjrzxqu-0x6u8nwy	5	2011-07-27
2	ieslbzqucldszsqm0ecsq	4	2012-06-14
3	g-wvgaisbqqamhlennbyoda	5	2010-05-27
4	lujfq2r5qfjg_6exmrcagw	5	2012-01-05

	text	type_x \
0	wife took birthday breakfast excellent weather...	review
1	idea people give bad reviews place goes show p...	review
2	love gyro plate rice good also dig candy selec...	review
3	rosie dakota love chaparral dog park convenien...	review
4	general manager scott petello good egg go deta...	review

	business_id	full_address	open	...	\
0	9ykzy9papeippoujetnvkg	6106 s 32nd st\nphoenix, az 85042	true	...	
1	zrjwvlyzejq1vaihdhyiow	4848 e chandler blvd\nphoenix, az 85044	true	...	
2	6orac4uyjcsj1lx0wzpvsa	1513 e apache blvd\ntempe, az 85281	true	...	
3	_lqqzuf4zzoyfcvxc0o6vg	5401 n hayden rd\nscottsdale, az 85250	true	...	
4	6ozyculrpktng2-lbrovtw	1357 s power road\nmesa, az 85206	true	...	

	word_count	word_density	punctuation_count	title_word_count \
0	155	5.698717948717949	21	16
1	257	5.213178294573644	36	25
2	16	4.470588235294118	3	2
3	76	5.441558441558442	18	13
4	86	5.390804597701149	38	12

	upper_case_word_count	stopword_count	line_count \
0	3	71	1
1	6	134	1
2	1	6	1
3	2	33	1
4	2	44	1

	txt_tokenized \
0	<filter object at 0x7fd8ff1bc410>

```

1 <filter object at 0x7fd8f1404fd0>
2 <filter object at 0x7fd8f14048d0>
3 <filter object at 0x7fd90275cf50>
4 <filter object at 0x7fd90047f990>

                                txt_stemmed \
0 [wife, took, birthday, breakfast, excel, weath...
1 [idea, peopl, give, bad, review, place, goe, s...
2 [love, gyro, plate, rice, good, also, dig, can...
3 [rosi, dakota, love, chaparr, dog, park, conve...
4 [gener, manag, scott, petello, good, egg, go, ...

                                txt_stemmed_sentence
0 wife took birthday breakfast excel weather per...
1 idea peopl give bad review place goe show plea...
2 love gyro plate rice good also dig candi select
3 rosi dakota love chaparr dog park conveni surr...
4 gener manag scott petello good egg go detail l...

[5 rows x 40 columns]

```

In [53]:

```

pos = df_merge3.loc[(df_merge3['score']=='positive')]
neg = df_merge3.loc[(df_merge3['score']=='negative')]

from wordcloud import WordCloud, ImageColorGenerator, STOPWORDS
stopwords = set(STOPWORDS)
pos_text = " ".join(review for review in pos.txt_stemmed_sentence)

# create figure
fig, ax = plt.subplots(figsize = (12,10))

wordcloud = WordCloud(width=1100, height=800, stopwords=stopwords).generate(pos_

plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")

ax.set_title('Positive Reviews', pad=15, fontsize = 20)
ax.title.set_position([.12, 0])

plt.show()

neg_text = " ".join(review for review in neg.txt_stemmed_sentence)

# create figure
fig, ax = plt.subplots(figsize = (12,10))

wordcloud = WordCloud(width=1100, height=800, stopwords=stopwords).generate(neg_

plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")

ax.set_title('Negative Reviews', pad=15, fontsize = 20)
ax.title.set_position([.13, 0])

plt.show()

```



localhost:8889/lab


```

all_words_proc = big_string_proc.split()
print(len(all_words_proc))

# create dictionary of word counts
fdist = FreqDist(all_words_proc)

# convert word counts to dataframe
fdist_df = pd.DataFrame(data=fdist.values(), index=fdist.keys(), columns=['word_c
fdist_df = fdist_df.sort_values('word_count', ascending=False)
top_25 = fdist_df.iloc[:25,:]
print(top_25)

# create labels and prettify the plot
plt.figure(figsize=(30,10))
plt.title('Top 25 Words', fontsize=36, pad=15)
plt.ylabel('Word Count', fontsize=30, labelpad=15)
plt.xticks(rotation=55, fontsize=28)
plt.yticks(fontsize=28)
plt.ylim(bottom=100, top=200000)

# plot top 25 words
plt.bar(top_25.index, top_25.word_count, color='purple')


# prepare to save and display
plt.tight_layout()
plt.show()

```

```

3349026
      word_count
place      36712
good       33369
food       30896
like       27648
great      25101
go         24018
get        23261
time       22616
one        21406
order      18553
realli     16821
servic     16142
would      15408
love       15145
tri        15031
back       14604
also       12522
dont       12408
even       11463
im         11438
nice       11425
make       11174
look       11051
littl      10887
restaur    10664

```

In [17]:

```

from sklearn.feature_selection import chi2
from sklearn.feature_extraction.text import TfidfVectorizer

# Check most important features in the dataset with chi2 score test

print('\nBelow are the most important features based on Chi2 score:\n')

Tfidf = TfidfVectorizer(min_df=5, ngram_range=(1, 2))
tfidf_features = Tfidf.fit_transform(df_merge3['txt_stemmed_sentence'])
tfidf_features.shape

N = 5
Number = 1
for rating in df_merge3['stars_x'].unique():
    features_chi2 = chi2(tfidf_features, df_merge3['stars_x'] == rating)
    indices = np.argsort(features_chi2[0])
    feature_names = np.array(Tfidf.get_feature_names()[indices])
    unigrams = [x for x in feature_names if len(x.split(' ')) == 1]
    bigrams = [x for x in feature_names if len(x.split(' ')) == 2]
    print("{} . {} :".format(Number, rating))
    print("\t Unigrams : \n\t. {}".format('\n\t. '.join(unigrams[-N:])))
    print("\t Bigrams : \n\t. {}".format('\n\t. '.join(bigrams[-N:])))
    Number += 1

print('\n')

```

Below are the most important features based on Chi2 score:

1. 5 :

```

Unigrams :
. highli
. awesom
. love
. best
. amaz
Bigrams :
. great food
. pretti good
. love love
. highli recommend
. love place

```

2. 4 :

```

Unigrams :
. ok
. horribl

```

```

. tasti
. worst
. good
Bigrams :
. happi hour
. realli enjoy
. good food
. realli good
. four star
3. 2 :
    Unigrams :
    . overpr
    . ok
    . meh
    . mediocr
    . bland
    Bigrams :
    . experienc better
    . want like
    . noth special
    . wont back
    . two star
4. 3 :
    Unigrams :
    . averag
    . pretti
    . aok
    . decent
    . ok
    Bigrams :
    . noth special
    . 35 star
    . food ok
    . pretti good
    . three star
5. 1 :
    Unigrams :
    . terribl
    . told
    . rude
    . horribl
    . worst
    Bigrams :
    . place suck
    . food poison
    . zero star
    . wast time
    . never go

```

In [24]:

```

#df_merge3.drop('text',axis='columns', inplace=True)
print(list(df_merge3.columns))

df_merge3.drop('char_count',axis='columns', inplace=True)
df_merge3.drop('word_count',axis='columns', inplace=True)
df_merge3.drop('word_density',axis='columns', inplace=True)
df_merge3.drop('punctuation_count',axis='columns', inplace=True)
df_merge3.drop('title_word_count',axis='columns', inplace=True)
df_merge3.drop('upper_case_word_count',axis='columns', inplace=True)
df_merge3.drop('stopword_count',axis='columns', inplace=True)
df_merge3.drop('line_count',axis='columns', inplace=True)

```

```
df_merge3.drop('txt_tokenized',axis='columns', inplace=True)
df_merge3.drop('txt_stemmed',axis='columns', inplace=True)
```

```
['votes_x', 'user_id', 'review_id', 'stars_x', 'date', 'text', 'type_x', 'business_id', 'full_address', 'open', 'categories', 'city', 'review_count_x', 'name_x', 'neighborhoods', 'longitude', 'state', 'stars_y', 'latitude', 'type_y', 'checkin_info', 'type_x', 'votes_y', 'name_y', 'average_stars', 'review_count_y', 'type_y', 'char_count', 'word_count', 'word_density', 'punctuation_count', 'title_word_count', 'upper_case_word_count', 'stopword_count', 'line_count', 'txt_tokenized', 'txt_stemmed', 'txt_stemmed_sentence']
```

In [25]:

```
##### MODEL BUILDING AND EVALUATIONS #####

# split the dataset into training and validation datasets
train_x, valid_x, train_y, valid_y = model_selection.train_test_split(df_merge3[

# label encode the target variable
encoder = preprocessing.LabelEncoder()
train_y = encoder.fit_transform(train_y)
valid_y = encoder.fit_transform(valid_y)

##### Create features from text #####

# create a count vectorizer object
count_vect = CountVectorizer(analyzer='word', token_pattern=r'\w{1,}')
count_vect.fit(df_merge3['txt_stemmed_sentence'])

# word level tf-idf
tfidf_vect = TfidfVectorizer(analyzer='word', token_pattern=r'\w{1,}', max_features=
tfidf_vect.fit(df_merge3['txt_stemmed_sentence'])
xtrain_tfidf = tfidf_vect.transform(train_x)
xvalid_tfidf = tfidf_vect.transform(valid_x)

# ngram level tf-idf
tfidf_vect_ngram = TfidfVectorizer(analyzer='word', token_pattern=r'\w{1,}', ngram
tfidf_vect_ngram.fit(df_merge3['txt_stemmed_sentence'])
xtrain_tfidf_ngram = tfidf_vect_ngram.transform(train_x)
xvalid_tfidf_ngram = tfidf_vect_ngram.transform(valid_x)

# characters level tf-idf
tfidf_vect_ngram_chars = TfidfVectorizer(analyzer='char', token_pattern=r'\w{1,}
tfidf_vect_ngram_chars.fit(df_merge3['txt_stemmed_sentence'])
xtrain_tfidf_ngram_chars = tfidf_vect_ngram_chars.transform(train_x)
xvalid_tfidf_ngram_chars = tfidf_vect_ngram_chars.transform(valid_x)

# transform the training and validation data using count vectorizer object
xtrain_count = count_vect.transform(train_x)
xvalid_count = count_vect.transform(valid_x)
```

In [23]:

```
##### Test with Truncated SVD technique to check the optimum number of components
# Dimensionality reduction

for num_components in range(100):
    svd = TruncatedSVD(n_components=num_components, random_state
```

```

X_svd = svd.fit_transform(xtrain_tfidf)
print(f"Total variance explained: {np.sum(svd.explained_vari

# Check explained variance ratio for 300 components/features
svd = TruncatedSVD(n_components=300, random_state=42)
X_svd = svd.fit_transform(xtrain_tfidf)
#print("\nTotal variance explained with 300 components: {np.sum(svd.explained_va
print(f"Total variance explained: {np.sum(svd.explained_variance_ratio_):.2f}",

```

```

Total variance explained: 0.00 number of components: 0
Total variance explained: 0.00 number of components: 1
Total variance explained: 0.01 number of components: 2
Total variance explained: 0.02 number of components: 3
Total variance explained: 0.02 number of components: 4
Total variance explained: 0.02 number of components: 5
Total variance explained: 0.03 number of components: 6
Total variance explained: 0.03 number of components: 7
Total variance explained: 0.04 number of components: 8
Total variance explained: 0.04 number of components: 9
Total variance explained: 0.04 number of components: 10
Total variance explained: 0.05 number of components: 11
Total variance explained: 0.05 number of components: 12
Total variance explained: 0.05 number of components: 13
Total variance explained: 0.05 number of components: 14
Total variance explained: 0.06 number of components: 15
Total variance explained: 0.06 number of components: 16
Total variance explained: 0.06 number of components: 17
Total variance explained: 0.06 number of components: 18
Total variance explained: 0.07 number of components: 19
Total variance explained: 0.07 number of components: 20
Total variance explained: 0.07 number of components: 21
Total variance explained: 0.07 number of components: 22
Total variance explained: 0.08 number of components: 23
Total variance explained: 0.08 number of components: 24
Total variance explained: 0.08 number of components: 25
Total variance explained: 0.08 number of components: 26
Total variance explained: 0.08 number of components: 27
Total variance explained: 0.09 number of components: 28
Total variance explained: 0.09 number of components: 29
Total variance explained: 0.09 number of components: 30
Total variance explained: 0.09 number of components: 31
Total variance explained: 0.09 number of components: 32
Total variance explained: 0.10 number of components: 33
Total variance explained: 0.10 number of components: 34
Total variance explained: 0.10 number of components: 35
Total variance explained: 0.10 number of components: 36
Total variance explained: 0.10 number of components: 37
Total variance explained: 0.10 number of components: 38
Total variance explained: 0.11 number of components: 39
Total variance explained: 0.11 number of components: 40
Total variance explained: 0.11 number of components: 41
Total variance explained: 0.11 number of components: 42
Total variance explained: 0.11 number of components: 43
Total variance explained: 0.11 number of components: 44
Total variance explained: 0.12 number of components: 45
Total variance explained: 0.12 number of components: 46
Total variance explained: 0.12 number of components: 47
Total variance explained: 0.12 number of components: 48
Total variance explained: 0.12 number of components: 49
Total variance explained: 0.12 number of components: 50
Total variance explained: 0.13 number of components: 51
Total variance explained: 0.13 number of components: 52
Total variance explained: 0.13 number of components: 53
Total variance explained: 0.13 number of components: 54

```

```

Total variance explained: 0.13 number of components: 55
Total variance explained: 0.13 number of components: 56
Total variance explained: 0.13 number of components: 57
Total variance explained: 0.14 number of components: 58
Total variance explained: 0.14 number of components: 59
Total variance explained: 0.14 number of components: 60
Total variance explained: 0.14 number of components: 61
Total variance explained: 0.14 number of components: 62
Total variance explained: 0.14 number of components: 63
Total variance explained: 0.14 number of components: 64
Total variance explained: 0.15 number of components: 65
Total variance explained: 0.15 number of components: 66
Total variance explained: 0.15 number of components: 67
Total variance explained: 0.15 number of components: 68
Total variance explained: 0.15 number of components: 69
Total variance explained: 0.15 number of components: 70
Total variance explained: 0.15 number of components: 71
Total variance explained: 0.16 number of components: 72
Total variance explained: 0.16 number of components: 73
Total variance explained: 0.16 number of components: 74
Total variance explained: 0.16 number of components: 75
Total variance explained: 0.16 number of components: 76
Total variance explained: 0.16 number of components: 77
Total variance explained: 0.16 number of components: 78
Total variance explained: 0.16 number of components: 79
Total variance explained: 0.17 number of components: 80
Total variance explained: 0.17 number of components: 81
Total variance explained: 0.17 number of components: 82
Total variance explained: 0.17 number of components: 83
Total variance explained: 0.17 number of components: 84
Total variance explained: 0.17 number of components: 85
Total variance explained: 0.17 number of components: 86
Total variance explained: 0.17 number of components: 87
Total variance explained: 0.18 number of components: 88
Total variance explained: 0.18 number of components: 89
Total variance explained: 0.18 number of components: 90
Total variance explained: 0.18 number of components: 91
Total variance explained: 0.18 number of components: 92
Total variance explained: 0.18 number of components: 93
Total variance explained: 0.18 number of components: 94
Total variance explained: 0.18 number of components: 95
Total variance explained: 0.18 number of components: 96
Total variance explained: 0.19 number of components: 97
Total variance explained: 0.19 number of components: 98
Total variance explained: 0.19 number of components: 99
Total variance explained: 0.36 number of components: 300

```

In [26]:

```

##### Train, Build and Evaluate the model #####
#Write a function to train the model classifier that will calculate the accuracy
#The function will also plot the confusion matrix and will print the classificat

def train_model(model_name,classifier, feature_vector_train, label, feature_vect
    # fit the training dataset on the classifier
    classifier.fit(feature_vector_train, label)

    # predict the labels on validation dataset
    predictions = classifier.predict(feature_vector_valid)

    if is_neural_net:
        predictions = predictions.argmax(axis=-1)
    cf_matrix=confusion_matrix(valid_y, predictions)
    print('\nConfusion matrix for' ,model_name,'is:\n')

```

```

print(confusion_matrix(valid_y, predictions))
print('\nClassification report for',model_name,'is:\n')
print(classification_report(valid_y, predictions,target_names=['1','2','3','4','5'],
index = ['0','1','2','3','4']
columns = ['1','2','3','4','5'])
cm_df = pd.DataFrame(cf_matrix,columns,index)
plt.figure(figsize=(5.5,4))
sns.heatmap(cm_df, annot=True,cmap="viridis" ,fmt='g')
plt.xticks([0,1,2,3,4])
plt.yticks([0,1,2,3,4])
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.title(model_name)
plt.show()

return metrics.accuracy_score(predictions, valid_y)

```

In [19]:

```

# Naive Bayes on Count Vectors
accuracy_nb_cv = train_model("Naive Bayes Count Vectors:",naive_bayes.Multinomi
print( "Accuracy of Naive Bayes, Count Vectors Model is: ", "{:.2%}".format(accuracy_nb_cv)

# Naive Bayes on Word Level TF IDF Vectors
accuracy_nb_tfidf = train_model("NB, WordLevel TF-IDF: ",naive_bayes.Multinomial
print( "Accuracy of Naive Bayes, WordLevel TF-IDF Model is: ", "{:.2%}".format(accuracy_nb_tfidf)

# Naive Bayes on Ngram Level TF IDF Vectors
accuracy_nb_ngtfidf = train_model("NB, N-Gram Vectors: ",naive_bayes.Multinomial
print( "Accuracy of Naive Bayes, N-Gram Vectors Model is: ", "{:.2%}".format(accuracy_nb_ngtfidf)

# Naive Bayes on Character Level TF IDF Vectors
accuracy_nb_ctfidf = train_model("NB, CharLevel Vectors: ",naive_bayes.Multinomi
print( "Accuracy of Naive Bayes, CharLevel Vectors Model is: ", "{:.2%}".format(accuracy_nb_ctfidf)

```

Confusion matrix for Naive Bayes Count Vectors: is:


```

[[ 927  326  261  236  108]
 [ 360  382  792  687  145]
 [ 162  145  957 2170  339]
 [ 151   39  439 5877 2155]
 [ 185   20   69 3147 4921]]

```

Classification report for Naive Bayes Count Vectors: is:

	precision	recall	f1-score	support
1	0.52	0.50	0.51	1858
2	0.42	0.16	0.23	2366
3	0.38	0.25	0.30	3773
4	0.49	0.68	0.57	8661
5	0.64	0.59	0.61	8342
accuracy			0.52	25000
macro avg	0.49	0.44	0.45	25000
weighted avg	0.52	0.52	0.51	25000



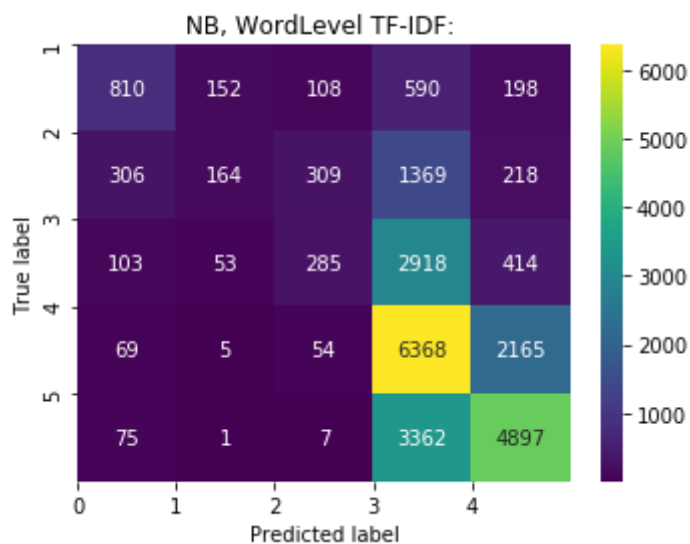
Accuracy of Naive Bayes, Count Vectors Model is: 52.26%

Confusion matrix for NB, WordLevel TF-IDF: is:

```
[[ 810  152  108  590  198]
 [ 306  164  309 1369  218]
 [ 103   53  285 2918  414]
 [   69    5   54 6368 2165]
 [   75    1    7 3362 4897]]
```

Classification report for NB, WordLevel TF-IDF: is:

	precision	recall	f1-score	support
1	0.59	0.44	0.50	1858
2	0.44	0.07	0.12	2366
3	0.37	0.08	0.13	3773
4	0.44	0.74	0.55	8661
5	0.62	0.59	0.60	8342
accuracy			0.50	25000
macro avg	0.49	0.38	0.38	25000
weighted avg	0.50	0.50	0.46	25000



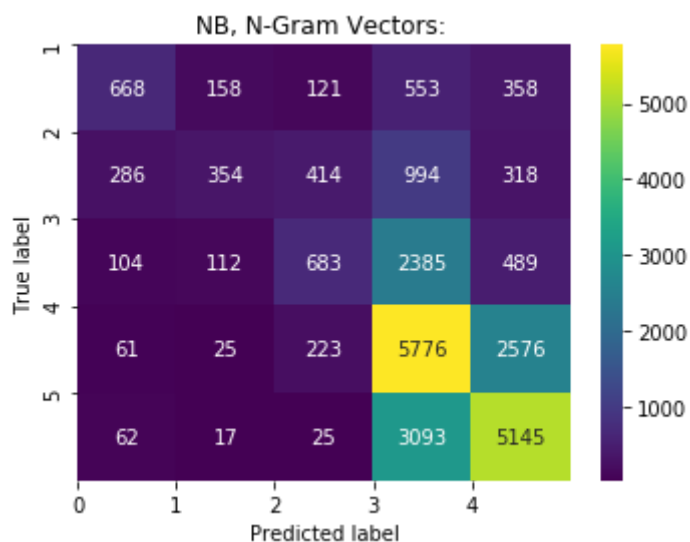
Accuracy of Naive Bayes, WordLevel TF-IDF Model is: 50.10%

Confusion matrix for NB, N-Gram Vectors: is:

```
[[ 668  158  121  553  358]
 [ 286  354  414  994  318]
 [ 104  112  683 2385  489]
 [  61   25  223 5776 2576]
 [  62   17   25 3093 5145]]
```

Classification report for NB, N-Gram Vectors: is:

	precision	recall	f1-score	support
1	0.57	0.36	0.44	1858
2	0.53	0.15	0.23	2366
3	0.47	0.18	0.26	3773
4	0.45	0.67	0.54	8661
5	0.58	0.62	0.60	8342
accuracy			0.51	25000
macro avg	0.52	0.39	0.41	25000
weighted avg	0.51	0.51	0.48	25000



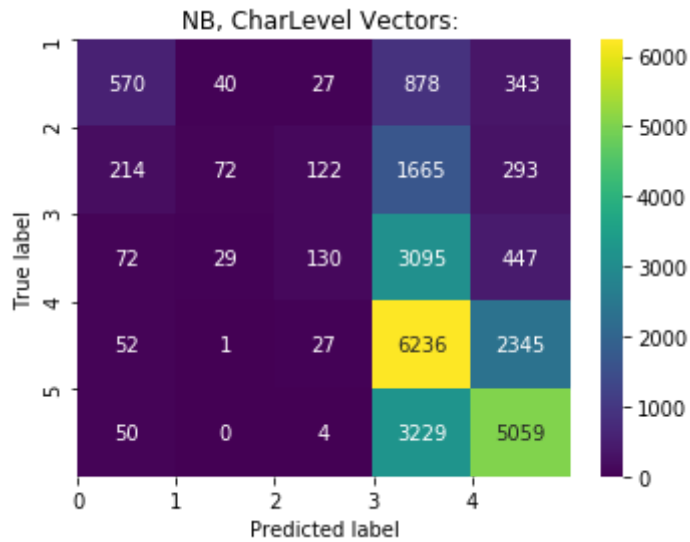
Accuracy of Naive Bayes, N-Gram Vectors Model is: 50.50%

Confusion matrix for NB, CharLevel Vectors: is:

```
[[ 570  40  27  878  343]
 [ 214  72 122 1665  293]
 [  72  29 130 3095  447]
 [  52   1  27 6236 2345]
 [  50   0   4 3229 5059]]
```

Classification report for NB, CharLevel Vectors: is:

	precision	recall	f1-score	support
1	0.59	0.31	0.40	1858
2	0.51	0.03	0.06	2366
3	0.42	0.03	0.06	3773
4	0.41	0.72	0.52	8661
5	0.60	0.61	0.60	8342
accuracy			0.48	25000
macro avg	0.51	0.34	0.33	25000
weighted avg	0.50	0.48	0.43	25000



Accuracy of Naive Bayes, CharLevel Vectors Model is: 48.27%

In [20]:

```
# Linear Classifier on Count Vectors
accuracy_lr_cv = train_model("LR, Count Vectors: ",linear_model.LogisticRegressi
print( "Accuracy of Liner Regression Count Vectors Model is: ", "{:.2%}".format(

# Linear Classifier on Word Level TF IDF Vectors
accuracy_lr_tfidf = train_model("LR, WordLevel TF-IDF: ",linear_model.LogisticRe
print( "Accuracy of Liner Regression TFIDF Model is: ", "{:.2%}".format(accuracy

# Linear Classifier on Ngram Level TF IDF Vectors
accuracy_lr_ngtfidf = train_model("LR, N-Gram Vectors: ",linear_model.LogisticRe
print( "Accuracy of Liner Regression N-Gram Level Model is: ", "{:.2%}".format(a


# Linear Classifier on Character Level TF IDF Vectors
accuracy_lr_ctfidf = train_model("LR, CharLevel TF-IDF: ",linear_model.LogisticR
print( "Accuracy of Liner Regression Char Level TFIDF Model is: ", "{:.2%}".form
```

Confusion matrix for LR, Count Vectors: is:

```
[[1032  333  133  210  150]
 [ 461  634  627  473  171]
 [ 144  365 1159 1614  491]
 [  72  142  754 4844 2849]
 [  60   45  177 2430 5630]]
```

Classification report for LR, Count Vectors: is:

	precision	recall	f1-score	support
1	0.58	0.56	0.57	1858
2	0.42	0.27	0.33	2366
3	0.41	0.31	0.35	3773
4	0.51	0.56	0.53	8661
5	0.61	0.67	0.64	8342
accuracy			0.53	25000
macro avg	0.50	0.47	0.48	25000
weighted avg	0.52	0.53	0.52	25000




Accuracy of Liner Regression Count Vectors Model is: 53.20%

Confusion matrix for LR, WordLevel TF-IDF: is:

```
[[1107  320  105  171  155]
 [ 433  621  622  490  200]
 [ 121  247 1180 1821  404]
 [  76   55  478 5316 2736]
 [  51   16   84 2290 5901]]
```

Classification report for LR, WordLevel TF-IDF: is:

	precision	recall	f1-score	support
1	0.62	0.60	0.61	1858
2	0.49	0.26	0.34	2366
3	0.48	0.31	0.38	3773
4	0.53	0.61	0.57	8661
5	0.63	0.71	0.67	8342
accuracy			0.56	25000
macro avg	0.55	0.50	0.51	25000
weighted avg	0.56	0.56	0.55	25000




Accuracy of Liner Regression TFIDF Model is: 56.50%

Confusion matrix for LR, N-Gram Vectors: is:

```
[[ 727  219  155  318  439]
 [ 299  528  485  660  394]
 [  98  198 1034 1761  682]
 [  77   69  466 5097 2952]
 [  79   36  117 2574 5536]]
```

Classification report for LR, N-Gram Vectors: is:

	precision	recall	f1-score	support
1	0.57	0.39	0.46	1858
2	0.50	0.22	0.31	2366
3	0.46	0.27	0.34	3773
4	0.49	0.59	0.53	8661
5	0.55	0.66	0.60	8342
accuracy			0.52	25000
macro avg	0.51	0.43	0.45	25000
weighted avg	0.51	0.52	0.50	25000




Accuracy of Liner Regression N-Gram Level Model is: 51.69%

Confusion matrix for LR, CharLevel TF-IDF: is:

```
[[ 999  265  137  214  243]
 [ 381  571  556  603  255]
 [ 128  227 1081 1871  466]
 [  93   56  387 5352 2773]
 [  70   18   69 2365 5820]]
```

Classification report for LR, CharLevel TF-IDF: is:

	precision	recall	f1-score	support
1	0.60	0.54	0.57	1858
2	0.50	0.24	0.33	2366
3	0.48	0.29	0.36	3773
4	0.51	0.62	0.56	8661
5	0.61	0.70	0.65	8342
accuracy			0.55	25000
macro avg	0.54	0.48	0.49	25000
weighted avg	0.55	0.55	0.54	25000



Accuracy of Liner Regression Char Level TFIDF Model is: 55.29%

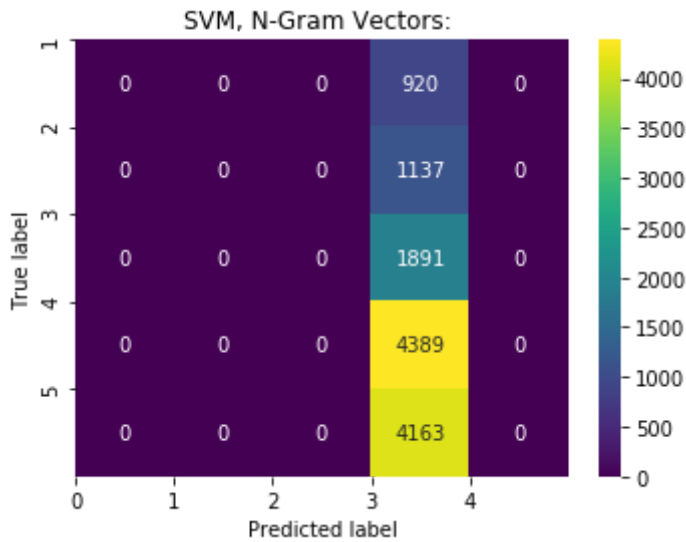
```
In [18]: # SVM on Ngram Level TF IDF Vectors
accuracy_svm_ngtfidf = train_model("SVM, N-Gram Vectors: ",svm.SVC(), xtrain_tfi
print( "Accuracy of SVM, N-Gram Vectors Model is: ", "{:.2%}".format(accuracy_sv
```

Confusion matrix for SVM, N-Gram Vectors: is:

[[0	0	0	920	0]
[0	0	0	1137	0]	
[0	0	0	1891	0]	
[0	0	0	4389	0]	
[0	0	0	4163	0]]	

Classification report for SVM, N-Gram Vectors: is:

	precision	recall	f1-score	support
1	0.00	0.00	0.00	920
2	0.00	0.00	0.00	1137
3	0.00	0.00	0.00	1891
4	0.35	1.00	0.52	4389
5	0.00	0.00	0.00	4163
accuracy			0.35	12500
macro avg	0.07	0.20	0.10	12500
weighted avg	0.12	0.35	0.18	12500



Accuracy of SVM, N-Gram Vectors Model is: 35.11%

In [19]:


```
# RF on Count Vectors
accuracy_rf_cv = train_model("RF, Count Vectors: ",ensemble.RandomForestClassifi
print( "Accuracy of Random Forest Count Vector Model is: ", "{:.2%}".format(accur
```

Confusion matrix for RF, Count Vectors: is:

```
[[ 245   95  131  264  185]
 [ 106  117  209  465  240]
 [   60  125  402  928  376]
 [   67  111  412 2323 1476]
 [   63   69  221 1655 2155]]
```

Classification report for RF, Count Vectors: is:

	precision	recall	f1-score	support
1	0.45	0.27	0.34	920
2	0.23	0.10	0.14	1137
3	0.29	0.21	0.25	1891
4	0.41	0.53	0.46	4389
5	0.49	0.52	0.50	4163
accuracy			0.42	12500
macro avg	0.37	0.33	0.34	12500
weighted avg	0.40	0.42	0.40	12500



Accuracy of Random Forest Count Vector Model is: 41.94%


```
In [20]: # RF on Word Level TF IDF Vectors
accuracy_rf_tfidf = train_model("RF, WordLevel TF-IDF: ",ensemble.RandomForestCl
print( "Accuracy of Random Forest word level Model is: ", "{:.2%}".format(accura
```

Confusion matrix for RF, WordLevel TF-IDF: is:

```
[[ 313  107  110  216  174]
 [ 139  147  220  435  196]
 [  82  122  432  936  319]
 [  75   91  355 2378 1490]
 [  48   56  167 1612 2280]]
```

Classification report for RF, WordLevel TF-IDF: is:

	precision	recall	f1-score	support
1	0.48	0.34	0.40	920
2	0.28	0.13	0.18	1137
3	0.34	0.23	0.27	1891
4	0.43	0.54	0.48	4389
5	0.51	0.55	0.53	4163
accuracy			0.44	12500
macro avg	0.41	0.36	0.37	12500
weighted avg	0.43	0.44	0.43	12500



Accuracy of Random Forest word level Model is: 44.40%


```
In [21]: # Extereme Gradient Boosting on Count Vectors
accuracy_xgb_cv = train_model("Xgb, Count Vectors: ",xgboost.XGBClassifier(), xt
print( "Accuracy of Xgradient boost count vector Model is: ", "{:.2%}".format(ac
```

Confusion matrix for Xgb, Count Vectors: is:

```
[[ 269   72   61  238  280]
 [   95  120  125  543  254]
 [   26   56  294 1147  368]
 [   14   13  106 2551 1705]
 [   10    5   15 1262 2871]]
```

Classification report for Xgb, Count Vectors: is:

	precision	recall	f1-score	support
1	0.65	0.29	0.40	920
2	0.45	0.11	0.17	1137
3	0.49	0.16	0.24	1891
4	0.44	0.58	0.50	4389
5	0.52	0.69	0.60	4163
accuracy			0.49	12500
macro avg	0.51	0.36	0.38	12500
weighted avg	0.49	0.49	0.46	12500



Accuracy of Xgradient boost count vector Model is: 48.84%


```
In [22]: # Exetereme Gradient Boosting on Word Level TF IDF Vectors
accuracy_xgb_tfidf = train_model("Xgb, WordLevel TF-IDF: ",xgboost.XGBClassifier
print( "Accuracy of Xgradient boost TFIDF vector Model is: ", "{:.2%}".format(ac
```

Confusion matrix for Xgb, WordLevel TF-IDF: is:

```
[[ 277   75   55  228  285]
 [   95  122  156  514  250]
 [   34   53  318 1137  349]
 [   16    9  105 2586 1673]
 [    9    2   18 1278 2856]]
```

Classification report for Xgb, WordLevel TF-IDF: is:

	precision	recall	f1-score	support
1	0.64	0.30	0.41	920
2	0.47	0.11	0.17	1137
3	0.49	0.17	0.25	1891
4	0.45	0.59	0.51	4389
5	0.53	0.69	0.60	4163
accuracy			0.49	12500
macro avg	0.52	0.37	0.39	12500
weighted avg	0.50	0.49	0.46	12500



Accuracy of Xgradient boost TFIDF vector Model is: 49.27%

In [23]:


```
# Exeterme Gradient Boosting on Character Level TF IDF Vectors
accuracy_xgb_cltfidf = train_model("Xgb, CharLevel Vectors: ",xgboost.XGBClassif
print( "Accuracy of Xgradient boost char level vector Model is:", "{:.2%}".forma
```

Confusion matrix for Xgb, CharLevel Vectors: is:

[278	61	54	251	276]
[113	121	157	523	223]
[37	60	358	1093	343]
[19	10	145	2570	1645]
[13	5	32	1284	2829]]

Classification report for Xgb, CharLevel Vectors: is:

	precision	recall	f1-score	support
1	0.60	0.30	0.40	920
2	0.47	0.11	0.17	1137
3	0.48	0.19	0.27	1891
4	0.45	0.59	0.51	4389
5	0.53	0.68	0.60	4163
accuracy			0.49	12500
macro avg	0.51	0.37	0.39	12500
weighted avg	0.49	0.49	0.46	12500



Accuracy of Xgradient boost char level vector Model is: 49.25%

In [37]:

```
train_x, valid_x, train_y, valid_y = model_selection.train_test_split(df_merge3[

# label encode the target variable
encoder = preprocessing.LabelEncoder()
train_y = encoder.fit_transform(train_y)
valid_y = encoder.fit_transform(valid_y)

# ngram level tf-idf
tfidf_vect_ngram = TfidfVectorizer(analyzer='word', token_pattern=r'\w{1,}', ngr
tfidf_vect_ngram.fit(df_merge3['txt_stemmed_sentence'])
xtrain_tfidf_ngram = tfidf_vect_ngram.transform(train_x).toarray()
xvalid_tfidf_ngram = tfidf_vect_ngram.transform(valid_x).toarray()

from keras.models import Sequential
from keras import layers

input_dim = xtrain_tfidf_ngram.shape[1] # Number of features

model = Sequential()
model.add(layers.Dense(10, input_dim=input_dim, activation='relu'))
model.add(layers.Dense(5, activation='sigmoid'))
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics
model.summary()

history = model.fit(xtrain_tfidf_ngram, train_y, epochs=100, verbose=False, vali

loss, accuracy = model.evaluate(xtrain_tfidf_ngram, train_y, verbose=False)
print("Training Accuracy: {:.4f}".format(accuracy))
loss, accuracy = model.evaluate(xvalid_tfidf_ngram, valid_y, verbose=False)
print("Testing Accuracy: {:.4f}".format(accuracy))
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_26 (Dense)	(None, 10)	50010
dense_27 (Dense)	(None, 5)	55
Total params: 50,065		
Trainable params: 50,065		
Non-trainable params: 0		
Training Accuracy: 0.8440		
Testing Accuracy: 0.4385		

In [35]:

```

# split the dataset into training and validation datasets
train_x, valid_x, train_y, valid_y = model_selection.train_test_split(df_merge3[

from sklearn.feature_extraction.text import CountVectorizer

vectorizer = CountVectorizer()
vectorizer.fit(df_merge3['txt_stemmed_sentence'])

X_train = vectorizer.transform(train_x)
X_test = vectorizer.transform(valid_x)

# label encode the target variable
encoder = preprocessing.LabelEncoder()
train_y = encoder.fit_transform(train_y)
valid_y = encoder.fit_transform(valid_y)

print(X_train.shape)
print(X_test.shape)
print(train_y.shape)
print(valid_y.shape)

def create_model_architecture(input_size):
    # create input layer
    input_layer = layers.Input((input_size, ), sparse=True)

    # create hidden layer
    hidden_layer = layers.Dense(100, activation="relu")(input_layer)

    # create output layer
    output_layer = layers.Dense(5, activation="softmax")(hidden_layer)

    classifier = models.Model(inputs = input_layer, outputs = output_layer)
    classifier.compile(optimizer=optimizers.Adam(), loss='binary_crossentropy')
    return classifier

classifier = create_model_architecture(X_train.shape[1])

```

```
accuracy_neural_network = train_model("Neural Network Ngram Level TF IDF Vector")
print("Accuracy of Neural Network Ngram Level TF IDF Vector Model is:", "{:.2%}"
```


```
(37500, 69182)
(12500, 69182)
(37500,)
(12500,)
1172/1172 [=====] - 27s 23ms/step - loss: 4.0633
```

Confusion matrix for Neural Network Ngram Level TF IDF Vector Model is:

```
[[ 14  13  23  33 833]
 [ 15  11   7  20 1105]
 [ 15  15  13  26 1819]
 [ 29  23  33  61 4184]
 [ 38  54  41  59 4016]]
```

Classification report for Neural Network Ngram Level TF IDF Vector Model is:

	precision	recall	f1-score	support
1	0.13	0.02	0.03	916
2	0.09	0.01	0.02	1158
3	0.11	0.01	0.01	1888
4	0.31	0.01	0.03	4330
5	0.34	0.95	0.50	4208
accuracy			0.33	12500
macro avg	0.19	0.20	0.12	12500
weighted avg	0.25	0.33	0.18	12500



Accuracy of Neural Network Ngram Level TF IDF Vector Model is: 32.92%

```
In [39]: train_x, valid_x, train_y, valid_y = model_selection.train_test_split(df_merge3[

# label encode the target variable
encoder = preprocessing.LabelEncoder()
train_y = encoder.fit_transform(train_y)
valid_y = encoder.fit_transform(valid_y)

# ngram level tf-idf
```

```
tfidf_vect_ngram = TfidfVectorizer(analyzer='word', token_pattern=r'\w{1,}', ngr
tfidf_vect_ngram.fit(df_merge3['txt_stemmed_sentence'])
xtrain_tfidf_ngram = tfidf_vect_ngram.transform(train_x).toarray()
xvalid_tfidf_ngram = tfidf_vect_ngram.transform(valid_x).toarray()
```

```
from keras.models import Sequential
from keras import layers
```

```
print(X_train.shape)
print(X_test.shape)
print(train_y.shape)
print(valid_y.shape)
```

```
def create_model_architecture(input_size):
    # create input layer
    input_layer = layers.Input((input_size, ), sparse=True)

    # create hidden layer
    hidden_layer = layers.Dense(10, activation="relu")(input_layer)

    # create output layer
    output_layer = layers.Dense(5, activation="softmax")(hidden_layer)

    classifier = models.Model(inputs = input_layer, outputs = output_layer)
    classifier.compile(optimizer=optimizers.Adam(), loss='sparse_categorical_crossentropy')
    return classifier
```

```
classifier = create_model_architecture(xtrain_tfidf_ngram.shape[1])
accuracy_neural_network_ngram = train_model("Neural Network Ngram Level TF IDF V
print( "Accuracy of Neural Network Ngram Level TF IDF Vector Model is:", "{:.2%}"
```

```
(37500, 69182)
(12500, 69182)
(37500,)
(12500,)
1172/1172 [=====] - 1s 712us/step - loss: 1.3172
```


Confusion matrix for Neural Network Ngram Level TF IDF Vector Model is:

```
[[ 53  257   89  304  271]
 [ 10  222  222  532  188]
 [  2   75  278 1232  356]
 [  1   19   85 2685 1489]
 [  4   18   17 1291 2800]]
```

Classification report for Neural Network Ngram Level TF IDF Vector Model is:

	precision	recall	f1-score	support
1	0.76	0.05	0.10	974
2	0.38	0.19	0.25	1174
3	0.40	0.14	0.21	1943
4	0.44	0.63	0.52	4279
5	0.55	0.68	0.61	4130

accuracy			0.48	12500
macro avg	0.51	0.34	0.34	12500
weighted avg	0.49	0.48	0.44	12500



Accuracy of Neural Network Ngram Level TF IDF Vector Model is: 48.30%

In [44]:

```
##### Word2vec word embedding using xgboost classifier #####

test_size = 0.3
random_state = 1234

X_train, X_test, y_train, y_test = model_selection.train_test_split(df_merge3['t
import numpy as np
from sklearn.base import BaseEstimator, TransformerMixin
from gensim.models import Word2Vec

class GensimWord2VecVectorizer(BaseEstimator, TransformerMixin):
    """
    Word vectors are averaged across to create the document-level vectors/featur
    gensim's own gensim.sklearn_api.W2VTransformer doesn't support out of vocabu
    hence we roll out our own.
    All the parameters are gensim.models.Word2Vec's parameters.
    https://radimrehurek.com/gensim/models/word2vec.html#gensim.models.word2vec.
    """

    def __init__(self, size=100, alpha=0.025, window=5, min_count=5, max_vocab_s
        sample=0.001, seed=1, workers=3, min_alpha=0.0001, sg=0, hs=0,
        ns_exponent=0.75, cbow_mean=1, hashfxn=hash, iter=5, null_word=
        trim_rule=None, sorted_vocab=1, batch_words=10000, compute_loss
        callbacks=(), max_final_vocab=None):
        self.size = size
        self.alpha = alpha
        self.window = window
        self.min_count = min_count
        self.max_vocab_size = max_vocab_size
        self.sample = sample
        self.seed = seed
        self.workers = workers
        self.min_alpha = min_alpha
        self.sg = sg
```

```

self.hs = hs
self.negative = negative
self.ns_exponent = ns_exponent
self.cbow_mean = cbow_mean
self.hashfxn = hashfxn
self.iter = iter
self.null_word = null_word
self.trim_rule = trim_rule
self.sorted_vocab = sorted_vocab
self.batch_words = batch_words
self.compute_loss = compute_loss
self.callbacks = callbacks
self.max_final_vocab = max_final_vocab

def fit(self, X, y=None):
    self.model_ = Word2Vec(
        sentences=X, corpus_file=None,
        size=self.size, alpha=self.alpha, window=self.window, min_count=self.
        max_vocab_size=self.max_vocab_size, sample=self.sample, seed=self.se
        workers=self.workers, min_alpha=self.min_alpha, sg=self.sg, hs=self.
        negative=self.negative, ns_exponent=self.ns_exponent, cbow_mean=self
        hashfxn=self.hashfxn, iter=self.iter, null_word=self.null_word,
        trim_rule=self.trim_rule, sorted_vocab=self.sorted_vocab, batch_word
        compute_loss=self.compute_loss, callbacks=self.callbacks,
        max_final_vocab=self.max_final_vocab)
    return self

def transform(self, X):
    X_embeddings = np.array([self._get_embedding(words) for words in X])
    return X_embeddings

def _get_embedding(self, words):
    valid_words = [word for word in words if word in self.model_.wv.vocab]
    if valid_words:
        embedding = np.zeros((len(valid_words), self.size), dtype=np.float32)
        for idx, word in enumerate(valid_words):
            embedding[idx] = self.model_.wv[word]

        return np.mean(embedding, axis=0)
    else:
        return np.zeros(self.size)

gensim_word2vec_tr = GensimWord2VecVectorizer(size=50, min_count=3, sg=1, alpha=
xgb = xgboost.XGBClassifier(learning_rate=0.01, n_estimators=100, n_jobs=-1)
w2v_xgb = Pipeline([('w2v', gensim_word2vec_tr), ('xgb', xgb)])
#w2v_xgb

w2v_xgb.fit(X_train, y_train)

y_train_pred = w2v_xgb.predict(X_train)
print('\nConfusion Matrix for word2vec xgboost model:\n')

```

```

y_test_pred = w2v_xgb.predict(X_test)
print(confusion_matrix(y_test, y_test_pred))
accuracy_w2v_xgb=metrics.accuracy_score(y_test, y_test_pred)
print( "Accuracy of XGBoost word2vec model classifier is:", "{:.2%}".format(metr
cf_matrix=confusion_matrix(y_test, y_test_pred)
index = ['0','1','2','3','4']
columns = ['0','1','2','3','4']
cm_df = pd.DataFrame(cf_matrix,columns,index)
plt.figure(figsize=(5.5,4))
sns.heatmap(cm_df, annot=True,cmap="viridis" ,fmt='g')
plt.xticks([0,1,2,3,4])
plt.yticks([0,1,2,3,4])
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.title("word2vec xgboost")
plt.show()

print('\n Classification report for word2vec xgboost model:\n')
print(classification_report(y_test, y_test_pred,target_names=['1','2','3','4','5']

```


Confusion Matrix for word2vec xgboost model:

```

[[ 2    0    0  756  383]
 [ 0    0    0  892  456]
 [ 0    0    0 1644  716]
 [ 1    0    0 3091 2027]
 [ 1    0    0 2363 2668]]

```

Accuracy of XGBoost word2vec model classifier is: 38.41%



Classification report for word2vec xgboost model:

	precision	recall	f1-score	support
1	0.50	0.00	0.00	1141
2	0.00	0.00	0.00	1348
3	0.00	0.00	0.00	2360
4	0.35	0.60	0.45	5119
5	0.43	0.53	0.47	5032
accuracy			0.38	15000
macro avg	0.26	0.23	0.18	15000
weighted avg	0.30	0.38	0.31	15000

In [27]: *# Function to Plot ROC curve for multi-class classification using different mode*

```
def plot_roc(model_name,clf,xtrain_count,y_train,xvalid_count,y_test):
    #clf = MultinomialNB(alpha=.01)
    clf.fit(xtrain_count, y_train)
    pred = clf.predict(xvalid_count)
    pred_prob = clf.predict_proba(xvalid_count)
    print('\n',model_name,'\n')
    # roc curve for classes

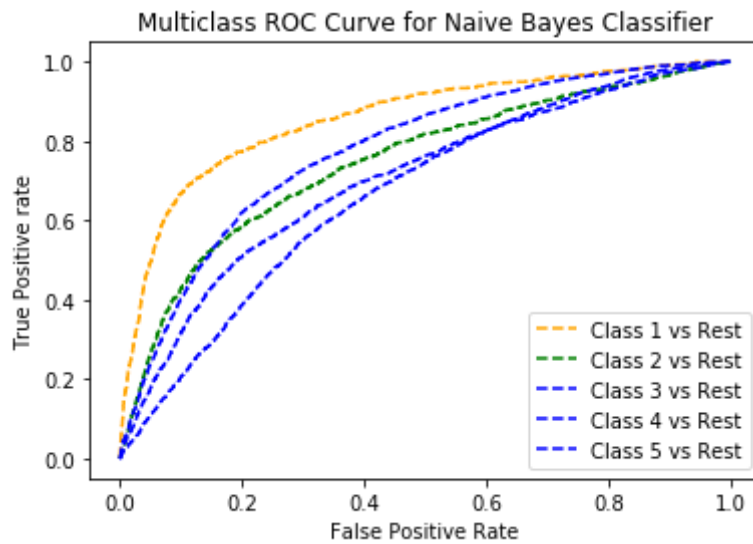
    fpr = {}
    tpr = {}
    thresh ={}

    n_class = 5
    for i in range(n_class):
        fpr[i], tpr[i], thresh[i] = roc_curve(y_test, pred_prob[:,i], pos_label=
        print('AUC for Class {}: {}'.format(i+1, auc(fpr[i], tpr[i])))
        # plotting
    plt.plot(fpr[0], tpr[0], linestyle='--',color='orange', label='Class 1 vs Re
    plt.plot(fpr[1], tpr[1], linestyle='--',color='green', label='Class 2 vs Res
    plt.plot(fpr[2], tpr[2], linestyle='--',color='blue', label='Class 3 vs Rest
    plt.plot(fpr[3], tpr[3], linestyle='--',color='blue', label='Class 4 vs Rest
    plt.plot(fpr[4], tpr[4], linestyle='--',color='blue', label='Class 5 vs Rest
    plt.title(model_name)
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive rate')
    plt.legend(loc='best')
    plt.show()

##### Plot ROC Curve for 4 different models #####
plot_roc("Multiclass ROC Curve for Naive Bayes Classifier",MultinomialNB(alpha=.
plot_roc("Multiclass ROC Curve for Linear Regression classifier",OneVsRestClassi
#plot_roc("Multiclass ROC curve for SVM classifier",OneVsRestClassifier(svm.SVC(
plot_roc("Multiclass ROC curve for Random Foreset Classifier ",RandomForestClass
```

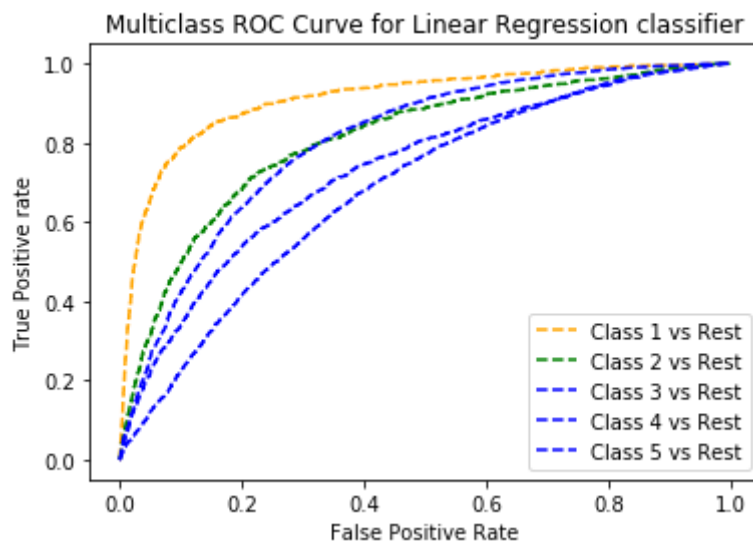
Multiclass ROC Curve for Naive Bayes Classifier

```
AUC for Class 1: 0.857274061205751
AUC for Class 2: 0.7469341678064164
AUC for Class 3: 0.7025143189626548
AUC for Class 4: 0.6697238429213672
AUC for Class 5: 0.777215664481072
```



Multiclass ROC Curve for Linear Regression classifier

AUC for Class 1: 0.9115248794062333
AUC for Class 2: 0.808149981612878
AUC for Class 3: 0.7325261115582609
AUC for Class 4: 0.6844827056833233
AUC for Class 5: 0.8036822386692285



Multiclass ROC curve for Random Foreset Classifier

AUC for Class 1: 0.7665607144464763
AUC for Class 2: 0.659734461937398
AUC for Class 3: 0.6210340460235904
AUC for Class 4: 0.5956354856607776
AUC for Class 5: 0.7077988247761994

